

# Next-POI Recommendation

*A TLMR-faithful baseline on Foursquare NYC*

*Architecture, training procedure & results*

PyTorch · PyTorch Geometric · Google Colab T4

Repository: [github.com/6ym6n/PFE\\_IMPLEMENTATION](https://github.com/6ym6n/PFE_IMPLEMENTATION)

Final report — May 2026

# Executive summary

This report documents an end-to-end implementation of a next-POI (Point-of-Interest) recommender on the Foursquare NYC check-in benchmark. The architecture is a two-layer GCN over a hybrid POI graph (co-visit  $\cup$  geographic kNN) feeding a GRU sequence model concatenated with a user embedding, followed by an MLP scoring head over the full POI vocabulary. The implementation faithfully follows the TLMR (Triple-Layered Modular Recommender) spec used as the project blueprint, with no "silent improvements."

**Result.** On the standard NYC test split (per-user chronological 70/10/20), the trained model reaches **HR@1 = 0.187**, **HR@5 = 0.476**, **HR@10 = 0.588**, **NDCG@10 = 0.375**, and **MRR = 0.316**. HR@1 lands at the top of the LSTM/STGCN tier reported in the LLM4POI benchmark table — exactly the band targeted for an honest baseline. No leakage symptoms (HR@1 > 0.20 with this architecture would be suspicious).

The remainder of the report walks through the problem formulation, the architecture stage by stage with diagrams, the training procedure, the ranking metrics, and a discussion of the results in context.

## 1. Problem formulation

Let  $U$  be the set of users and  $V$  the vocabulary of POIs. Each user  $u \in U$  produces a chronological sequence of check-ins; we segment these into **sessions** by splitting at temporal gaps larger than 24 hours (Foursquare convention). A session of length  $L$  is an ordered tuple  $S = (p_1, \dots, p_L)$  with associated context features  $\Delta d_t$  (haversine distance from  $p_{t-1}$  to  $p_t$  in km) and  $\Delta t_t$  (time gap from  $p_{t-1}$  to  $p_t$  in hours).

The **next-POI prediction** task is: given a prefix  $(p_1, \dots, p_T)$  from a session and the user identity  $u$ , predict a probability distribution over  $V$  for the next POI  $p_{T+1}$ . Quality is measured by ranking metrics (HR@k, NDCG@k, MRR) computed over the entire POI vocabulary — no negative sampling, no candidate restriction.

$$P(p_{T+1} = v \mid S, u) = \text{softmax}(\hat{y})_v, \text{ where } \hat{y} = f_\theta(S, u; G) \in \mathbb{R}^{|V|}$$

## 2. Architecture overview

The model is a **three-stage pipeline**: (1) feature extraction lifts each input into a vector representation, (2) a sequence model integrates the prefix into a single context vector, and (3) a scoring head produces logits over all POIs.




Figure 1 — End-to-end forward pipeline. Inputs (top): the POI prefix, per-step context scalars, and the user id. Each is encoded independently (middle row), POI features and context are concatenated per step, the sequence is consumed by a GRU, the final hidden state is concatenated with the user embedding, and an MLP head produces logits over the entire vocabulary.

### 3. Stage 1 — feature extraction

#### 3.1 Hybrid POI graph

POIs are not isolated tokens — they sit in geographic space and are linked by user behavior. The model exploits both signals by building a single hybrid graph  $G = (V, E_{cov} \cup E_{geo})$ :

- **Co-visit edges ( $E_{cov}$ ).** For every pair  $(p_i, p_j)$ , count how many distinct users transitioned  $p_i \rightarrow p_j$  as consecutive check-ins within a session, in **training data only**. Add the undirected edge if the count  $\geq \tau_{cov} = 3$ . Building this from train+val+test would be subtle leakage.
- **Geographic kNN edges ( $E_{geo}$ ).** Each POI is connected to its  $k = 10$  nearest neighbors by haversine distance, computed with a BallTree. This captures spatial proximity even for POIs with sparse co-visit history.

The two edge sets are unioned (de-duplicated) and symmetrized into a PyG edge\_index tensor of shape  $(2, 2|E|)$ . On NYC the resulting graph has 5,284 co-visit + 31,790 kNN edges, union 34,898, with degree min/median/max = 10/13/79 and **no isolated nodes** — every POI receives graph signal.




Figure 2 — Hybrid POI graph (illustrative subset). Thin blue edges are geographic kNN; thick orange edges are training-derived co-visit transitions. Co-visit edges connect POIs that users actually traveled between, even if they are not geographically close.

### 3.2 GCN encoder

A learnable embedding table  $E^{(0)} \in \mathbb{R}^{|V| \times d_p}$  ( $d_p = 128$ , Xavier init) holds initial POI embeddings. Two GCN layers propagate them through the hybrid graph using the standard symmetric normalization:

$$E^{(\ell+1)} = \sigma(D^{-1/2} \tilde{A} D^{-1/2} E^{(\ell)} W^{(\ell)}), \ell = 0, 1$$

where  $\tilde{A} = A + I$  adds self-loops and  $D$  is the corresponding degree matrix. ReLU + dropout( $p=0.2$ ) is applied between the two layers. The output  $h_p^{\text{GCN}} = E^{(2)} \in \mathbb{R}^{|V| \times d_p}$  is the matrix of POI features used downstream — the sequence model gathers rows from this matrix by POI index. Two layers is the standard depth: more leads to over-smoothing (POI features collapse toward each other).




Figure 3 — Two-layer GCN over the hybrid graph. The full vocabulary is propagated once per forward pass (not per example), then indexed by the input POI ids.

### 3.3 Context encoder

Per-step context is the haversine distance  $\Delta d_t$  and the time gap  $\Delta t_t$  from the previous POI in the session. Two parallel MLPs lift each scalar into  $\mathbb{R}^{16}$ , then the halves are concatenated:

$$c_t = [ \text{MLP}_d(\Delta d_t) ; \text{MLP}_t(\Delta t_t) ] \in \mathbb{R}^{d_c=32}$$

Each MLP is  $\text{Linear}(1 \rightarrow 16) \rightarrow \text{ReLU} \rightarrow \text{Linear}(16 \rightarrow 16)$ . At session boundaries the previous POI is undefined, so  $\Delta d_1 = \Delta t_1 = 0$  by convention.  $\Delta d$  is clipped to  $[0, 100]$  km and  $\Delta t$  to  $[0, 24]$  hours, which prevents extreme outliers (e.g. a typo in a timestamp) from destabilizing the MLPs.




Figure 4 — Context encoder. Two independent MLPs let the model learn separate non-linear mappings for spatial and temporal gaps before they are combined.

### 3.4 User embedding

Each user has a learnable embedding  $e_u \in \mathbb{R}^{(d_u=64)}$  stored in an `nn.Embedding(|U|, 64)` table (Xavier init). This captures stable, user-specific preferences (e.g. a user who systematically prefers cafés to bars) that are not present in the session prefix itself.

## 4. Stage 2 — sequence model

### 4.1 GRU encoder

Per-step input is the concatenation of POI feature and context:  $x_t = [h_{\{p_t\}^{GCN}}; c_t] \in \mathbb{R}^{(d_p + d_c)} = \mathbb{R}^{160}$ . Sessions in a minibatch have variable length, so they are right-padded to the batch maximum and consumed by the GRU via `pack_padded_sequence` (lengths must live on CPU for PyTorch's packing). The final hidden state  $h_T \in \mathbb{R}^{(d_h=128)}$  summarizes the entire prefix.




Figure 5 — GRU unrolled over a session prefix. Only the final hidden state  $h_T$  is propagated downstream; intermediate states are discarded.

### 4.2 Scoring head

The user embedding is concatenated with the GRU summary,  $z = [h_T; e_u] \in \mathbb{R}^{(d_h + d_u)} = \mathbb{R}^{192}$ , and fed to a two-layer MLP with hidden dim 256:

$$\hat{y} = W_2 \cdot \text{ReLU}(W_1 z + b_1) + b_2 \in \mathbb{R}^{|V|}$$

Output dimension is the full POI vocabulary (5,135 on NYC after filtering) so this is the parameter-heaviest layer. Dropout( $p=0.2$ ) is applied after the ReLU. The softmax is implicit in the cross-entropy loss; for ranking we use the raw logits.

## 5. Training procedure

Each session of length  $L$  is converted into  $L-1$  supervised examples by taking every prefix (length  $1..L-1$ )  $\rightarrow$  next-POI pair. Histories are right-padded per batch and consumed by the model in a single pass. The loss is plain cross-entropy over the  $|V|$ -way logits:

$$L = - (1/N) \cdot \sum_i \log P(y_i | S_i, u_i)$$




Figure 6 — Training pipeline. The GCN runs once per forward pass on the entire vocabulary (not per example), then per-step POI features are gathered by index — this is critical for throughput.

### 5.1 Hyperparameters (locked from spec)

| Hyperparameter                  | Value                                                        |
|---------------------------------|--------------------------------------------------------------|
| POI embedding dim $d_p$         | 128                                                          |
| User embedding dim $d_u$        | 64                                                           |
| Context dim $d_c$               | 32 (16 + 16)                                                 |
| GRU hidden dim $d_h$            | 128                                                          |
| MLP head hidden dim             | 256                                                          |
| GCN layers                      | 2                                                            |
| Dropout                         | 0.2                                                          |
| Optimizer                       | Adam (lr=1e-3, weight_decay=1e-5)                            |
| Batch size                      | 64                                                           |
| Max epochs / patience           | 50 / 8 (early stop on val HR@10)                             |
| Gradient clip (L2 norm)         | 5.0                                                          |
| Co-visit threshold $\tau_{cov}$ | 3 (training data only)                                       |
| Geographic kNN k                | 10 (haversine)                                               |
| Filter                          | $\geq 10$ check-ins/user, $\geq 10$ visitors/POI (iterative) |
| Session split gap               | 24 h                                                         |

Train / val / test

70 / 10 / 20 chronological per user

Random seed

42

## 6. Evaluation metrics

Given the logits  $\hat{y} \in \mathbb{R}^{|V|}$  and the true next POI  $y$ , the **rank** of the target is the number of POIs scored strictly higher than the target, plus one:

$$\text{rank}(y) = |\{v \in V : \hat{y}_v > \hat{y}_y\}| + 1$$

All three metrics derive from this rank, averaged over every test example:

### 6.1 Hit Rate at k (HR@k)

$$\text{HR@k} = (1/N) \cdot \sum_i [\text{rank}(y_i) \leq k]$$

Fraction of examples where the true next POI appears in the top-k recommendations. HR@1 is the strictest — it equals the exact-match accuracy.

### 6.2 Normalized Discounted Cumulative Gain (NDCG@k)

$$\text{NDCG@k} = (1/N) \cdot \sum_i [\text{rank}(y_i) \leq k] / \log_2(\text{rank}(y_i) + 1)$$

Like HR@k, but rewards higher ranks more — a target at rank 1 contributes 1, at rank 2 contributes  $1/\log_2(3) \approx 0.63$ , and so on. Standard ranking metric in IR and recommender systems.

### 6.3 Mean Reciprocal Rank (MRR)

$$\text{MRR} = (1/N) \cdot \sum_i 1 / \text{rank}(y_i)$$

Average of the reciprocal of the rank — captures the full ranking quality without a cutoff. A model that ranks the target second on every example would score MRR = 0.5; one that always ranks it first would score 1.0.

All metrics are computed over the **full POI vocabulary** (no negative sampling, no candidate restriction). This matches the GETNext / STHGCN evaluation protocol used by the LLM4POI benchmark we compare against — a critical detail since candidate-restricted evaluation can inflate metrics by 2–3×.

## 7. Results on NYC

### 7.1 Data pipeline sanity

Each preprocessing stage was instrumented; here are the actual numbers from the run:

| Stage                                        | Check-ins | Users           | POIs           |
|----------------------------------------------|-----------|-----------------|----------------|
| Raw download                                 | 227,428   | 1,083           | 38,333         |
| After iterative filter ( $\geq 10/\geq 10$ ) | 147,938   | 1,083           | 5,135          |
| After sessionization ( $\geq 2$ per session) | —         | 23,867 sessions | avg length 5.2 |
| Train split (70%)                            | 91,148    | 16,232 sessions | —              |
| Val split (10%)                              | 8,825     | 1,909 sessions  | —              |

|                  |        |                |   |
|------------------|--------|----------------|---|
| Test split (20%) | 24,734 | 5,726 sessions | — |
|------------------|--------|----------------|---|

Hybrid graph: 5,284 co-visit edges + 31,790 kNN edges, union 34,898 unique edges. Degree distribution min=10, median=13, mean=13.6, max=79. Zero isolated POIs (every node receives graph signal).

Materialized examples after generating prefix→target pairs: **74,916 train, 6,916 val, 19,008 test**.

## 7.2 Test metrics

| Metric  | Value  | Interpretation                                             |
|---------|--------|------------------------------------------------------------|
| HR@1    | 0.1874 | true next POI ranked #1 on 18.7% of test cases             |
| HR@5    | 0.4760 | true next POI in top 5 on 47.6% of test cases              |
| HR@10   | 0.5879 | true next POI in top 10 on 58.8% of test cases             |
| NDCG@5  | 0.3386 | rank-discounted hit rate at k=5                            |
| NDCG@10 | 0.3751 | rank-discounted hit rate at k=10                           |
| MRR     | 0.3163 | average reciprocal rank $\approx 1/3.16 \approx$ rank 3.16 |

## 7.3 Comparison to literature

Numbers below are HR@1 / Acc@1 on Foursquare NYC reproduced from Table 3 of the LLM4POI paper. Our model is bolded.

| Model                       | HR@1 (NYC)   | Family                                |
|-----------------------------|--------------|---------------------------------------|
| LSTM                        | 0.130        | Pure sequence model                   |
| STGCN                       | 0.180        | Spatio-temporal GCN                   |
| <b>This baseline (ours)</b> | <b>0.187</b> | <b>Hybrid GCN + GRU + user (TLMR)</b> |
| STAN                        | 0.220        | Self-attention spatio-temporal        |
| GETNext                     | 0.240        | Transformer + trajectory flow         |
| STHGCN                      | 0.270        | Hyper-graph spatio-temporal           |
| LLM4POI                     | 0.340        | Frozen LLM with prompt engineering    |

**Reading the result.** HR@1 = 0.187 sits between STGCN (0.180) and STAN (0.220) — i.e. at the very top of the "LSTM/STGCN tier" targeted as an honest baseline. The spec explicitly warns that HR@1 above 0.20 with this architecture is "suspicious — re-check the chronological split and graph construction for leakage." We are below that threshold, validating that the chronological split, the training-only co-visit edges, and the per-user prefix generation are all clean.

# 8. Discussion

## 8.1 What worked

- **The hybrid graph structure was useful.** Most POIs land near the kNN-only floor of 10 neighbors, but high-traffic POIs accumulate up to 79 neighbors via the co-visit edges. The fact that no POI was



isolated means the GCN never silently fails to propagate.

- **Per-step concat (not gating).** Concatenating context with the GCN POI feature is the simplest possible fusion and worked well at this scale. Path B in the project roadmap proposes moving  $\Delta_d, \Delta_t$  into ST-GRU gates — that is the next obvious experiment but was deliberately not done here.
- **User embedding contributed.** Foursquare users are a stable population; per-user preferences over POIs are strong enough that a 64-dim lookup adds measurable signal on top of the session prefix.

## 8.2 What we are not claiming

- We are **not** claiming SOTA. The baseline lands  $\approx 8$  HR@1 points below GETNext and  $\approx 15$  below LLM4POI. Both rely on considerably more sophisticated mechanisms (attention over trajectory flow graphs, frozen LLM in-context reasoning) that are out of scope here.
- We are **not** claiming the architecture is novel. The TLMR three-stage shape (graph features + sequence + user) is standard in the POI literature; the value of this work is a clean, reproducible reference implementation honest about its tier.

## 8.3 Honest caveats

- **Single seed.** Results are reported for seed 42 only. POI-rec results in the literature have known sensitivity to initialization; reporting mean  $\pm$  std across 3+ seeds is planned for the thesis (Week 9 of the roadmap).
- **NYC only in this report.** The original spec also covers TKY; that run is intentionally deferred. The pipeline is dataset-agnostic so adding TKY is a one-cell change.
- **No ablations yet.** The next planned experiment removes each component (GCN, user embedding, context) one at a time to quantify its contribution. Until that is done we cannot say *which* piece is doing the work — only that the assembly lands in the expected band.

## 8.4 Path B — beyond the baseline

Four high-leverage extensions are listed in the project roadmap, ordered by expected impact:

- Replace the MLP-to- $|V|$  head with inner-product scoring  $\hat{y}_v = (h_T + e_u)^T h_v^{GCN}$ . Same TLMR shape, much better transfer of graph signal.
- Move  $\Delta_d, \Delta_t$  into ST-GRU gates (STGN-style) instead of side concat. Typically gains 3–5 HR points.
- Add hour-of-day and day-of-week cyclic features at every step. Typically gains 2–4 HR points on Foursquare.
- Weight the co-visit edges by transition frequency (DLIR-style adaptive adjacency) instead of unweighted.

# 9. Phase 2 — itinerary recommendation

A second phase extends the next-POI baseline to **itinerary recommendation**: given a query (start POI, fixed end POI, and a length budget  $K$ ), produce an **ordered route** of  $K$  POIs. Quality is measured by **pairs-F1** (Chen 2016) — F1 over correctly-ordered POI pairs — on the length $\geq 3$  test subset ( $n=2,880$ ), where ordering actually matters (length-2 sessions are trivially perfect).

## 9.1 Two strategies

- **Strategy A — frozen rollout.** Decode the *frozen* next-POI model as an itinerary generator: iterate it from the start, mask visited POIs (loop-free), reserve the fixed end for the last hop. No retraining.
- **Strategy B — pointer network (trained).** Reuse the GCN + user encoder; replace the MLP head with an inner-product pointer; train end-to-end on whole trajectories (teacher forcing), early-stop on val pairs-F1. B-v1 omits  $\Delta d/\Delta t$  context; **B-v2** adds it back into the decoder input.

## 9.2 Results (NYC, length $\geq 3$ test, n=2,880)

| Method                      | pairs-F1       | set-F1 | exact | vs floor  |
|-----------------------------|----------------|--------|-------|-----------|
| A — frozen rollout (greedy) | 0.2887         | 0.609  | 0.054 | — (floor) |
| A — frozen rollout (beam 3) | 0.2902         | 0.610  | 0.057 | +0.0015   |
| B-v1 — pointer, no context  | 0.2585         | 0.578  | 0.043 | -0.0302   |
| B-v2 — pointer + context    | (run on Colab) | —      | —     | —         |

**Reading the result.** Beam barely beats greedy in Strategy A (+0.5%) — the next-POI scorer is myopic, so better *decoding* cannot fix it. More importantly, the from-scratch trained pointer (B-v1) **under-performs the frozen baseline by 0.030 pairs-F1 (-10.5%)**. This is a genuine negative result (B-v1 trained cleanly: loss 8.01 $\rightarrow$ 3.41, val pairs-F1 peaking at 0.271).

**Why the frozen model wins.** Its next-POI engine was trained on  $\sim 75,000$  prefix $\rightarrow$ next examples (every prefix, all lengths), whereas B-v1 saw only 10,281 whole-trajectory examples and used leaner per-step inputs (no  $\Delta d/\Delta t$  context). The honest lesson: *naively training a dedicated itinerary model does not automatically beat cleverly decoding a strong next-POI model*. B-v2 (context-aware pointer) is the implemented next test of whether the missing  $\Delta d/\Delta t$  features close the gap.

## 10. Phase 2b — Flickr datasets (literature-comparable pairs-F1)

The NYC itinerary numbers above (pairs-F1  $\approx 0.26$ – $0.29$ ) are **not comparable** to the published trip-recommendation literature, which reports pairs-F1  $\approx 0.6$ – $0.8$  — but on a *different* benchmark (the small Flickr photo-trajectory datasets) under a *different* protocol. **Strategy D** re-runs the itinerary task on those exact datasets, under the exact published protocol, so the numbers land on the same scale as the papers. The low NYC numbers turn out to be a property of that data and protocol — not a bug.

**Data.** We use the canonical traj-{City}.csv / poi-{City}.csv files — Chen 2016's own preprocessing output (mirrored in the tour-cikm16 and DeepTrip repos) — so our trajectories are *identical* to the literature's, removing all preprocessing ambiguity. Five cities: Toronto, Osaka, Glasgow, Edinburgh, Melbourne (27–88 POIs each). Evaluation is on the length $\geq 3$  subset (335 / 47 / 112 / 634 / 442 trajectories respectively).

**Protocol** (matched to the papers exactly): leave-one-trajectory-out cross-validation, refitting each fold on the other length $\geq 3$  trajectories; the query gives the first and last POI plus the length  $K$ ; length $\geq 3$ , loop-free trajectories only; metrics are point-F1 and order-aware pairs-F1, reusing the same unit-tested implementation as Phase 2 (pinned against Chen 2016's reference calc\_pairsF1).

### 10.1 Protocol-faithfulness check (the honesty proof)

A **Random** baseline has no modelling choices, so it can only match Chen 2016's Random if the data, protocol, and metric are identical to the paper's. They match within noise — as does PoiPopularity:

| pairs-F1                  | Toronto | Osaka | Glasgow | Edinburgh | Melbourne |
|---------------------------|---------|-------|---------|-----------|-----------|
| Random — ours             | 0.298   | 0.301 | 0.301   | 0.270     | 0.227     |
| Random — Chen 2016        | 0.310   | 0.304 | 0.320   | 0.261     | 0.248     |
| PoiPopularity — ours      | 0.443   | 0.413 | 0.510   | 0.439     | 0.320     |
| PoiPopularity — Chen 2016 | 0.384   | 0.365 | 0.507   | 0.436     | 0.316     |

## 10.2 Our measured results (classical, leave-one-out, length $\geq 3$ )

| pairs-F1          | Toronto | Osaka | Glasgow | Edinburgh | Melbourne |
|-------------------|---------|-------|---------|-----------|-----------|
| Random            | 0.298   | 0.301 | 0.301   | 0.270     | 0.227     |
| PoiPopularity     | 0.443   | 0.413 | 0.510   | 0.439     | 0.320     |
| Markov (greedy)   | 0.504   | 0.421 | 0.587   | 0.449     | 0.333     |
| MarkovPath (beam) | 0.528   | 0.398 | 0.543   | 0.452     | 0.346     |

**The headline.** Our pairs-F1 is **0.23–0.59** — **squarely on the published 0.26–0.85 scale**, versus 0.26–0.29 on Foursquare NYC. The thesis goal (numbers on the literature's scale) is met by the classical baselines alone; the learned GCN + pointer model is measured in §10.5. (MarkovPath beam-searches for the maximum-transition-likelihood path — a surrogate for, not identical to, pairs-F1 — so it slightly trails greedy Markov on Osaka.)

## 10.3 Published comparison (directly-comparable subset)

Same protocol (leave-one-out, endpoints given, length $\geq 3$ ), 0–1 scale, curated in `src/flickr/published.py`:

| pairs-F1 (published)     | Toronto | Osaka | Glasgow | Edinburgh |
|--------------------------|---------|-------|---------|-----------|
| PoiRank (Chen 2016)      | 0.518   | 0.511 | 0.548   | 0.432     |
| Rank+Markov (Chen 2016)  | 0.512   | 0.486 | 0.545   | 0.444     |
| DeepTrip (2019)          | 0.748*  | 0.755 | 0.782   | 0.660     |
| SelfTrip (2022)          | 0.835   | 0.851 | 0.818   | 0.779     |
| CTLTR (via AR-Trip 2024) | 0.748   | 0.719 | 0.763   | 0.681     |
| AR-Trip (2024)           | 0.839   | 0.828 | 0.820   | 0.808     |

There is no single SOTA across all cities on pairs-F1: **AR-Trip leads on Toronto / Glasgow / Edinburgh, but SelfTrip is best on Osaka (0.851 > 0.828)**. \*DeepTrip's paper reports only Edinburgh/Glasgow/Osaka; the Toronto value is from SelfTrip's reproduction. Melbourne is reported almost only by the classical line (best there: Rank+Markov 0.351).

## 10.4 Honest deviations

- **Markov  $\neq$  Chen's Markov.** Ours is a raw empirical first-order transition matrix (Laplace-smoothed); Chen's is feature-factored. Ours therefore differs and on the data-rich cities *exceeds* Chen's published Markov — a modelling difference, not a protocol bug. Random and PoiPopularity are the clean reproductions.
- **PoiRank / Rank+Markov are cited, not re-implemented** (they need a per-query rankSVM); their numbers come straight from Chen 2016.
- **Excluded as not comparable:** POIBERT (80/20 chronological split, percentage-scale F1, no pairs-F1), DLIR 2025 (8-hour session split, different F1  $\approx 0.49$ ), TourEmbedding (no pairs-F1, percentage-scale F1, query gives only the first POI not first+last, no leave-one-out).
- **Learned model.** A light self-contained GCN + GRU pointer (src/flickr/pointer.py, no PyTorch-Geometric needed) is trained per leave-one-out fold on Colab via colab\_flickr.ipynb — measured result in §10.5.

## 10.5 Measured learned model (Colab, GCN + pointer, GPU)

A real run of colab\_flickr.ipynb (60 epochs per leave-one-out fold, T4 GPU):

| pairs-F1         | Toronto | Osaka | Glasgow | Edinburgh | Melbourne |
|------------------|---------|-------|---------|-----------|-----------|
| Pointer (beam 3) | 0.431   | 0.399 | 0.489   | 0.414     | 0.312     |
| Pointer (greedy) | 0.430   | 0.362 | 0.486   | 0.414     | 0.309     |
| point-F1 (beam)  | 0.705   | 0.684 | 0.734   | 0.689     | 0.614     |
| ref: Random      | 0.298   | 0.301 | 0.301   | 0.271     | 0.227     |
| ref: Markov      | 0.513   | 0.421 | 0.592   | 0.450     | 0.333     |

**Honest reading (the Strategy-D lesson).** The light from-scratch pointer reaches pairs-F1 **0.31–0.49** — **on the published scale and clearly above Random**, and its point-F1 ( $\sim 0.68$ – $0.73$ ) shows it recovers the *right POIs*. But its **ordering trails the simple Markov baseline on every city**, and it is far below the neural SOTA band (0.66–0.85). Exactly as on Foursquare NYC (Strategy B), *naively training a dedicated pointer on the tiny per-fold data does not beat a strong simple model* — the published neural methods reach SOTA with self-supervised / contrastive pre-training, trajectory augmentation, and adversarial training that this light model omits.

Implemented levers to close the gap (opt-in in PointerConfig, reported side-by-side in the notebook): a decode-time **Markov transition prior** (blends the fold's empirical  $\log P(j|i)$  into the pointer logits, directly targeting the weak-ordering gap), an optional **user embedding**, and **early stopping** on an internal validation split.

## Appendix A — Implementation summary

The codebase is split into a thin src/ package and a Colab-ready notebook (train\_poi.ipynb) that imports from it. Module map:

| Module                 | Public API                                                                                    |
|------------------------|-----------------------------------------------------------------------------------------------|
| src/data/preprocess.py | load_foursquare, iterative_filter, reindex, haversine_km, build_sessions, chronological_split |
| src/data/graph.py      | build_covisit_edges, build_knn_edges, build_hybrid_graph                                      |

|                          |                                                                                                                      |
|--------------------------|----------------------------------------------------------------------------------------------------------------------|
| src/data/dataset.py      | POISessionDataset, collate_fn                                                                                        |
| src/models/components.py | POIGraphEncoder, ContextEncoder                                                                                      |
| src/models/next_poi.py   | NextPOIModel                                                                                                         |
| src/train.py             | make_loaders, train_one_epoch, train_model                                                                           |
| src/eval.py              | evaluate_model, fmt_metrics                                                                                          |
| src/itinerary/           | query, decode, eval_itinerary (pairs_f1/set_f1), pointer_model, run_itinerary (Phase 2)                              |
| src/flickr/              | data (load_city), evaluate (evaluate_loocv), baselines, pointer (FlickrPointerNet), published, run_flickr (Phase 2b) |

Tests: 74 pytest cases — Section-6 smoke test, controlled-rank evaluation, padding/dtype and shape tests for every module, the itinerary pairs-F1 / decode-invariant suite, and the Flickr suite (loader, trajectory construction, pairs-F1 pinned against Chen 2016's reference, leave-one-out splitter, baseline validity) — all pass on a fresh CPU-only environment.

## Appendix B — Software stack

| Component         | Version                  |
|-------------------|--------------------------|
| Python            | 3.10+ (Colab default)    |
| PyTorch           | 2.3.0                    |
| PyTorch Geometric | 2.5.3                    |
| NumPy             | 1.26.4                   |
| pandas            | 2.1.3                    |
| scikit-learn      | 1.4.2                    |
| pyarrow           | 15.0.2 (parquet I/O)     |
| tqdm              | 4.66.4                   |
| pytest            | 8.1.1 (test runner only) |
| Hardware          | Google Colab T4 (16 GB)  |

## Appendix C — Reproducibility checklist

- Single global seed (42) propagated to random, numpy, torch, and CUDA RNG.
- All preprocessed artifacts (parquets, edge\_index, meta.json) persisted to Google Drive at /MyDrive/poi-rec/data/processed/NYC/ so the training cell can be re-run without redoing preprocessing.
- Best-by-val-HR@10 checkpoint and per-epoch metric history written to disk after every epoch — sessions can be resumed from checkpoints/NYC/latest.pt after a Colab disconnect.
- Public repository: [github.com/6ym6n/PFE\\_IMPLEMENTATION](https://github.com/6ym6n/PFE_IMPLEMENTATION) with pinned requirements.txt, pytest suite, and this report.