

Next-POI Recommendation

A TLMR-faithful baseline on Foursquare NYC

Architecture, training procedure & results

PyTorch · PyTorch Geometric · Google Colab T4

Repository: github.com/6ym6n/PFE_IMPLEMENTATION

Final report — May 2026

Executive summary

This report documents an end-to-end implementation of a next-POI (Point-of-Interest) recommender on the Foursquare NYC check-in benchmark. The architecture is a two-layer GCN over a hybrid POI graph (co-visit $\cup$ geographic kNN) feeding a GRU sequence model concatenated with a user embedding, followed by an MLP scoring head over the full POI vocabulary. The implementation faithfully follows the TLMR (Triple-Layered Modular Recommender) spec used as the project blueprint, with no "silent improvements."

Result. On the standard NYC test split (per-user chronological 70/10/20), the trained model reaches **HR@1 = 0.187**, **HR@5 = 0.476**, **HR@10 = 0.588**, **NDCG@10 = 0.375**, and **MRR = 0.316**. HR@1 lands at the top of the LSTM/STGCN tier reported in the LLM4POI benchmark table — exactly the band targeted for an honest baseline. No leakage symptoms (HR@1 > 0.20 with this architecture would be suspicious).

The remainder of the report walks through the problem formulation, the architecture stage by stage with diagrams, the training procedure, the ranking metrics, and a discussion of the results in context.

1. Problem formulation

Let U be the set of users and V the vocabulary of POIs. Each user $u \in U$ produces a chronological sequence of check-ins; we segment these into **sessions** by splitting at temporal gaps larger than 24 hours (Foursquare convention). A session of length L is an ordered tuple $S = (p_1, \dots, p_L)$ with associated context features Δd_t (haversine distance from p_{t-1} to p_t in km) and Δt_t (time gap from p_{t-1} to p_t in hours).

The **next-POI prediction** task is: given a prefix $(p_1, \dots, p_T)$ from a session and the user identity u , predict a probability distribution over V for the next POI p_{T+1} . Quality is measured by ranking metrics (HR@k, NDCG@k, MRR) computed over the entire POI vocabulary — no negative sampling, no candidate restriction.

$$P(p_{T+1} = v \mid S, u) = \text{softmax}(\mathbf{z})_v, \text{ where } \mathbf{z} = f_\theta(S, u; G) \in \mathbb{R}^{|V|}$$

2. Architecture overview

The model is a **three-stage pipeline**: (1) feature extraction lifts each input into a vector representation, (2) a sequence model integrates the prefix into a single context vector, and (3) a scoring head produces logits over all POIs.




Figure 1 — End-to-end forward pipeline. Inputs (top): the POI prefix, per-step context scalars, and the user id. Each is encoded independently (middle row), POI features and context are concatenated per step, the sequence is consumed by a GRU, the final hidden state is concatenated with the user embedding, and an MLP head produces logits over the entire vocabulary.

3. Stage 1 — feature extraction

3.1 Hybrid POI graph

POIs are not isolated tokens — they sit in geographic space and are linked by user behavior. The model exploits both signals by building a single hybrid graph $G = (V, E_{cov} \cup E_{geo})$:

- **Co-visit edges (E_{cov}).** For every pair (p_i, p_j) , count how many distinct users transitioned $p_i \rightarrow p_j$ as consecutive check-ins within a session, in **training data only**. Add the undirected edge if the count $\geq \tau_{cov} = 3$. Building this from train+val+test would be subtle leakage.
- **Geographic kNN edges (E_{geo}).** Each POI is connected to its $k = 10$ nearest neighbors by haversine distance, computed with a BallTree. This captures spatial proximity even for POIs with sparse co-visit history.

The two edge sets are unioned (de-duplicated) and symmetrized into a PyG `edge_index` tensor of shape $(2, 2|E|)$. On NYC the resulting graph has 5,284 co-visit + 31,790 kNN edges, union 34,898, with degree min/median/max = 10/13/79 and **no isolated nodes** — every POI receives graph signal.




Figure 2 — Hybrid POI graph (illustrative subset). Thin blue edges are geographic kNN; thick orange edges are training-derived co-visit transitions. Co-visit edges connect POIs that users actually traveled between, even if they are not geographically close.

3.2 GCN encoder

A learnable embedding table $E \in \mathbb{R}^{|V| \times d_p}$ ($d_p = 128$, Xavier init) holds initial POI embeddings. Two GCN layers propagate them through the hybrid graph using the standard symmetric normalization:

$$E^{(l+1)} = \sigma(D^{-1/2} \tilde{A} D^{-1/2} E^{(l)} W^{(l)}), \quad l = 0, 1$$

where $\tilde{A} = A + I$ adds self-loops and D is the corresponding degree matrix. ReLU + dropout($p=0.2$) is applied between the two layers. The output $h_p^{\text{GCN}} = E^{(2)} \in \mathbb{R}^{|V| \times d_p}$ is the matrix of POI features used downstream — the sequence model gathers rows from this matrix by POI index. Two layers is the standard depth: more leads to over-smoothing (POI features collapse toward each other).




Figure 3 — Two-layer GCN over the hybrid graph. The full vocabulary is propagated once per forward pass (not per example), then indexed by the input POI ids.

3.3 Context encoder

Per-step context is the haversine distance Δd_t and the time gap Δt_t from the previous POI in the session. Two parallel MLPs lift each scalar into $\mathbb{R}^{16}$, then the halves are concatenated:

$$c_t = [MLP_d(\Delta d_t) ; MLP_t(\Delta t_t)] \in \mathbb{R}^{(d_c=32)}$$

Each MLP is $\text{Linear}(1 \rightarrow 16) \rightarrow \text{ReLU} \rightarrow \text{Linear}(16 \rightarrow 16)$. At session boundaries the previous POI is undefined, so $\Delta d_1 = \Delta t_1 = 0$ by convention. Δd is clipped to $[0, 100]$ km and Δt to $[0, 24]$ hours, which prevents extreme outliers (e.g. a typo in a timestamp) from destabilizing the MLPs.




Figure 4 — Context encoder. Two independent MLPs let the model learn separate non-linear mappings for spatial and temporal gaps before they are combined.

3.4 User embedding

Each user has a learnable embedding $e_u \in \mathbb{R}^{(d_u=64)}$ stored in an `nn.Embedding(|U|, 64)` table (Xavier init). This captures stable, user-specific preferences (e.g. a user who systematically prefers cafés to bars) that are not present in the session prefix itself.

4. Stage 2 — sequence model

4.1 GRU encoder

Per-step input is the concatenation of POI feature and context: $x_t = [h_{\{p_t\}}^{GCN} ; c_t] \in \mathbb{R}^{(d_p + d_c) = \mathbb{R}^{160}}$. Sessions in a minibatch have variable length, so they are right-padded to the batch maximum and consumed by the GRU via `pack_padded_sequence` (lengths must live on CPU for PyTorch's packing). The final hidden state $h_T \in \mathbb{R}^{(d_h=128)}$ summarizes the entire prefix.




Figure 5 — GRU unrolled over a session prefix. Only the final hidden state h_T is propagated downstream; intermediate states are discarded.

4.2 Scoring head

The user embedding is concatenated with the GRU summary, $z = [h_T ; e_u] \in \mathbb{R}^{(d_h + d_u) = \mathbb{R}^{192}}$, and fed to a two-layer MLP with hidden dim 256:

$$\mathbf{y} = W \cdot \text{ReLU}(Wz + b) + b \in \mathbb{R}^{|V|}$$

Output dimension is the full POI vocabulary (5,135 on NYC after filtering) so this is the parameter-heaviest layer. Dropout($p=0.2$) is applied after the ReLU. The softmax is implicit in the cross-entropy loss; for ranking we use the raw logits.

5. Training procedure

Each session of length L is converted into $L-1$ supervised examples by taking every prefix (length $1..L-1$) $\rightarrow$ next-POI pair. Histories are right-padded per batch and consumed by the model in a single pass. The loss is plain cross-entropy over the $|V|$ -way logits:

$$L = - (1/N) \cdot \sum \log P(y | S, u)$$




Figure 6 — Training pipeline. The GCN runs once per forward pass on the entire vocabulary (not per example), then per-step POI features are gathered by index — this is critical for throughput.

5.1 Hyperparameters (locked from spec)

Hyperparameter	Value
POI embedding dim d_p	128
User embedding dim d_u	64
Context dim d_c	32 (16 + 16)
GRU hidden dim d_h	128
MLP head hidden dim	256
GCN layers	2
Dropout	0.2
Optimizer	Adam (lr=1e-3, weight_decay=1e-5)
Batch size	64
Max epochs / patience	50 / 8 (early stop on val HR@10)
Gradient clip (L2 norm)	5.0
Co-visit threshold τ_{cov}	3 (training data only)
Geographic kNN k	10 (haversine)
Filter	≥ 10 check-ins/user, ≥ 10 visitors/POI (iterative)

Session split gap	24 h
Train / val / test	70 / 10 / 20 chronological per user
Random seed	42

6. Evaluation metrics

Given the logits $\mathbf{v} \in \mathbb{R}^{|V|}$ and the true next POI y , the **rank** of the target is the number of POIs scored strictly higher than the target, plus one:

$$\text{rank}(y) = |\{v \in V : \mathbf{v}_y > \mathbf{v}_y\}| + 1$$

All three metrics derive from this rank, averaged over every test example:

6.1 Hit Rate at k (HR@k)

$$\text{HR@k} = (1/N) \cdot \sum_{\mathbf{v}} [\text{rank}(y_{\mathbf{v}}) \leq k]$$

Fraction of examples where the true next POI appears in the top-k recommendations. HR@1 is the strictest — it equals the exact-match accuracy.

6.2 Normalized Discounted Cumulative Gain (NDCG@k)

$$\text{NDCG@k} = (1/N) \cdot \sum_{\mathbf{v}} [\text{rank}(y_{\mathbf{v}}) \leq k] / \log(\text{rank}(y_{\mathbf{v}}) + 1)$$

Like HR@k, but rewards higher ranks more — a target at rank 1 contributes 1, at rank 2 contributes $1/\log(3) \approx 0.63$, and so on. Standard ranking metric in IR and recommender systems.

6.3 Mean Reciprocal Rank (MRR)

$$\text{MRR} = (1/N) \cdot \sum_{\mathbf{v}} 1 / \text{rank}(y_{\mathbf{v}})$$

Average of the reciprocal of the rank — captures the full ranking quality without a cutoff. A model that ranks the target second on every example would score $\text{MRR} = 0.5$; one that always ranks it first would score 1.0.

All metrics are computed over the **full POI vocabulary** (no negative sampling, no candidate restriction). This matches the GETNext / STHGCN evaluation protocol used by the LLM4POI benchmark we compare against — a critical detail since candidate-restricted evaluation can inflate metrics by 2–3×.

7. Results on NYC

7.1 Data pipeline sanity

Each preprocessing stage was instrumented; here are the actual numbers from the run:

Stage	Check-ins	Users	POIs
Raw download	227,428	1,083	38,333
After iterative filter ($\geq 10/\geq 10$)	147,938	1,083	5,135
After sessionization (≥ 2 per session)	—	23,867 sessions	avg length 5.2

Train split (70%)	91,148	16,232 sessions	—
Val split (10%)	8,825	1,909 sessions	—
Test split (20%)	24,734	5,726 sessions	—

Hybrid graph: 5,284 co-visit edges + 31,790 kNN edges, union 34,898 unique edges. Degree distribution min=10, median=13, mean=13.6, max=79. Zero isolated POIs (every node receives graph signal).

Materialized examples after generating prefix→target pairs: **74,916 train, 6,916 val, 19,008 test**.

7.2 Test metrics

Metric	Value	Interpretation
HR@1	0.1874	true next POI ranked #1 on 18.7% of test cases
HR@5	0.4760	true next POI in top 5 on 47.6% of test cases
HR@10	0.5879	true next POI in top 10 on 58.8% of test cases
NDCG@5	0.3386	rank-discounted hit rate at k=5
NDCG@10	0.3751	rank-discounted hit rate at k=10
MRR	0.3163	average reciprocal rank $\approx 1/3.16 \approx$ rank 3.16

7.3 Comparison to literature

Numbers below are HR@1 / Acc@1 on Foursquare NYC reproduced from Table 3 of the LLM4POI paper. Our model is bolded.

Model	HR@1 (NYC)	Family
LSTM	0.130	Pure sequence model
STGCN	0.180	Spatio-temporal GCN
This baseline (ours)	0.187	Hybrid GCN + GRU + user (TLMR)
STAN	0.220	Self-attention spatio-temporal
GETNext	0.240	Transformer + trajectory flow
STHGCN	0.270	Hyper-graph spatio-temporal
LLM4POI	0.340	Frozen LLM with prompt engineering

Reading the result. HR@1 = 0.187 sits between STGCN (0.180) and STAN (0.220) — i.e. at the very top of the "LSTM/STGCN tier" targeted as an honest baseline. The spec explicitly warns that HR@1 above 0.20 with this architecture is "suspicious — re-check the chronological split and graph construction for leakage." We are below that threshold, validating that the chronological split, the training-only co-visit edges, and the per-user prefix generation are all clean.

8. Discussion

8.1 What worked

- **The hybrid graph structure was useful.** Most POIs land near the kNN-only floor of 10 neighbors, but high-traffic POIs accumulate up to 79 neighbors via the co-visit edges. The fact that no POI was isolated means the GCN never silently fails to propagate.
- **Per-step concat (not gating).** Concatenating context with the GCN POI feature is the simplest possible fusion and worked well at this scale. Path B in the project roadmap proposes moving Δd , Δt into ST-GRU gates — that is the next obvious experiment but was deliberately not done here.
- **User embedding contributed.** Foursquare users are a stable population; per-user preferences over POIs are strong enough that a 64-dim lookup adds measurable signal on top of the session prefix.

8.2 What we are not claiming

- We are **not** claiming SOTA. The baseline lands ≈ 8 HR@1 points below GETNext and ≈ 15 below LLM4POI. Both rely on considerably more sophisticated mechanisms (attention over trajectory flow graphs, frozen LLM in-context reasoning) that are out of scope here.
- We are **not** claiming the architecture is novel. The TLMR three-stage shape (graph features + sequence + user) is standard in the POI literature; the value of this work is a clean, reproducible reference implementation honest about its tier.

8.3 Honest caveats

- **Single seed.** Results are reported for seed 42 only. POI-rec results in the literature have known sensitivity to initialization; reporting mean $\pm$ std across 3+ seeds is planned for the thesis (Week 9 of the roadmap).
- **NYC only in this report.** The original spec also covers TKY; that run is intentionally deferred. The pipeline is dataset-agnostic so adding TKY is a one-cell change.
- **No ablations yet.** The next planned experiment removes each component (GCN, user embedding, context) one at a time to quantify its contribution. Until that is done we cannot say *which* piece is doing the work — only that the assembly lands in the expected band.

8.4 Path B — beyond the baseline

Four high-leverage extensions are listed in the project roadmap, ordered by expected impact:

- Replace the MLP-to- $|V|$ head with inner-product scoring $\mathbf{h}_T + \mathbf{e}_u \cdot \mathbf{h}_v^{GCN}$. Same TLMR shape, much better transfer of graph signal.
- Move Δd , Δt into ST-GRU gates (STGN-style) instead of side concat. Typically gains 3–5 HR points.
- Add hour-of-day and day-of-week cyclic features at every step. Typically gains 2–4 HR points on Foursquare.
- Weight the co-visit edges by transition frequency (DLIR-style adaptive adjacency) instead of unweighted.

Appendix A — Implementation summary

The codebase is split into a thin `src/` package and a Colab-ready notebook (`train_poi.ipynb`) that imports from it. Module map:

Module	Public API
<code>src/data/preprocess.py</code>	<code>load_foursquare</code> , <code>iterative_filter</code> , <code>reindex</code> , <code>haversine_km</code> , <code>build_sessions</code> , <code>chronological_split</code>
<code>src/data/graph.py</code>	<code>build_covisit_edges</code> , <code>build_knn_edges</code> , <code>build_hybrid_graph</code>
<code>src/data/dataset.py</code>	<code>POISessionDataset</code> , <code>collate_fn</code>
<code>src/models/components.py</code>	<code>POIGraphEncoder</code> , <code>ContextEncoder</code>
<code>src/models/next_poi.py</code>	<code>NextPOIModel</code>
<code>src/train.py</code>	<code>make_loaders</code> , <code>train_one_epoch</code> , <code>train_model</code>
<code>src/eval.py</code>	<code>evaluate_model</code> , <code>fnt_metrics</code>

Tests: 24 pytest cases including a Section-6 smoke test, a controlled-rank evaluation test, padding/dtype tests, and shape tests for every module — all pass on a fresh CPU-only environment.

Appendix B — Software stack

Component	Version
Python	3.10+ (Colab default)
PyTorch	2.3.0
PyTorch Geometric	2.5.3
NumPy	1.26.4
pandas	2.1.3
scikit-learn	1.4.2
pyarrow	15.0.2 (parquet I/O)
tqdm	4.66.4
pytest	8.1.1 (test runner only)
Hardware	Google Colab T4 (16 GB)

Appendix C — Reproducibility checklist

- Single global seed (42) propagated to random, numpy, torch, and CUDA RNG.
- All preprocessed artifacts (parquets, `edge_index`, `meta.json`) persisted to Google Drive at `/MyDrive/poi-rec/data/processed/NYC/` so the training cell can be re-run without redoing preprocessing.
- Best-by-val-HR@10 checkpoint and per-epoch metric history written to disk after every epoch — sessions can be resumed from `checkpoints/NYC/latest.pt` after a Colab disconnect.

- Public repository: github.com/6ym6n/PFE_IMPLEMENTATION with pinned requirements.txt, pytest suite, and this report.